

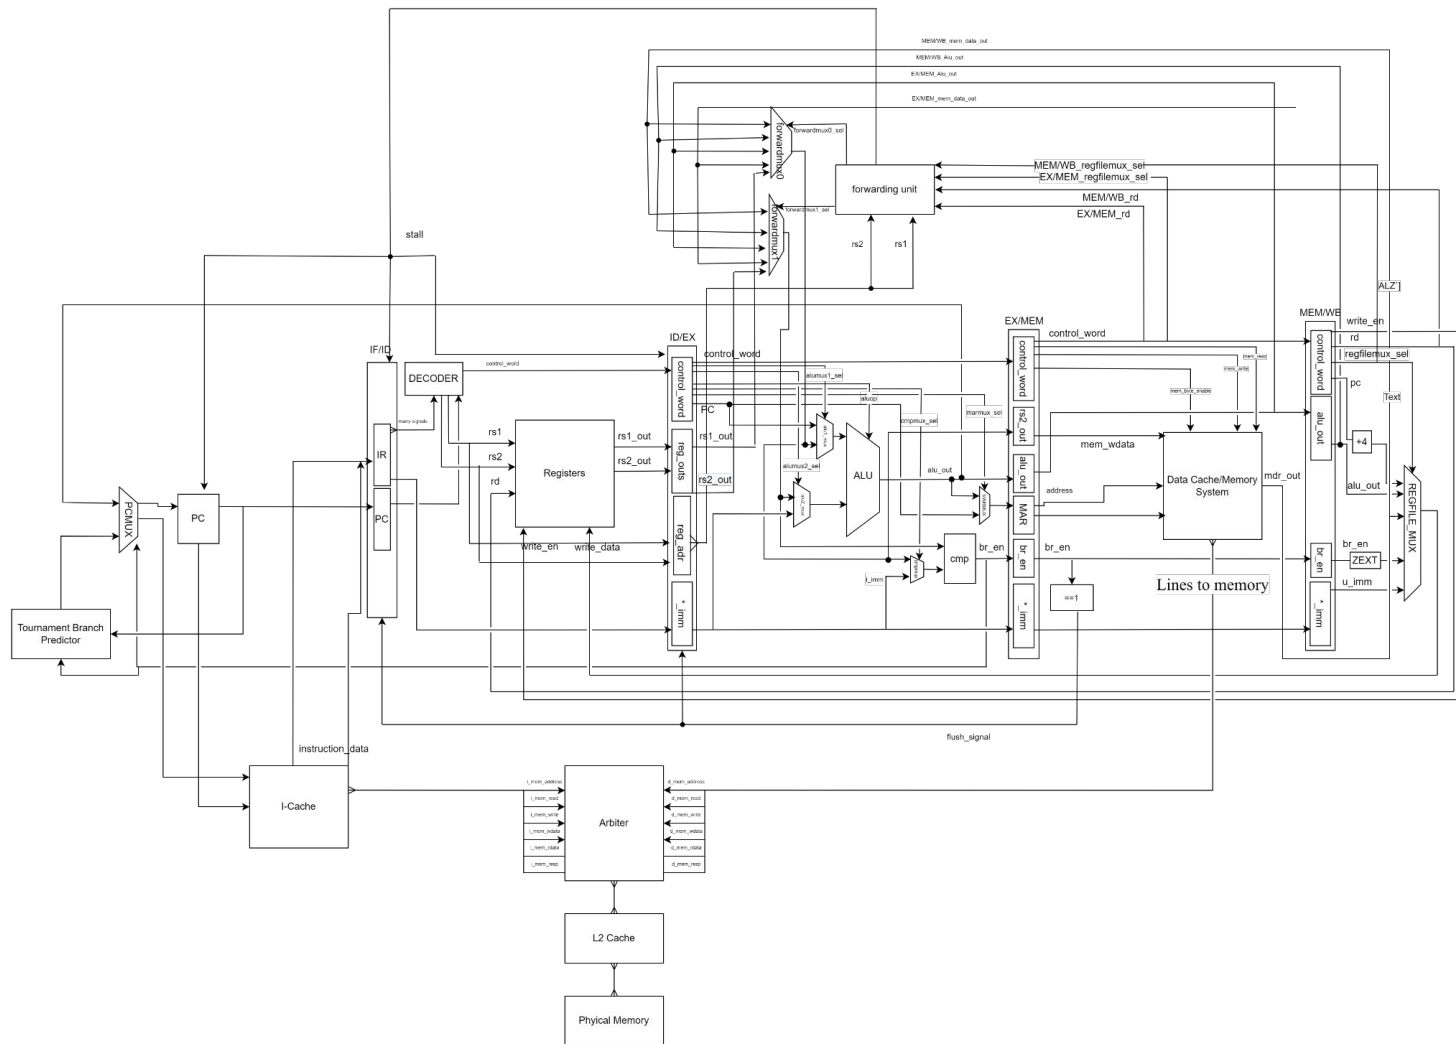
Snapdragon 411

Presented by Daniel Marks, Jason Liu, and Jincheng Liu

Overview

Implemented components

- (1) Basic 5-stages pipelined datapath. (cp 1)
- (2) Hazard detection & forwarding for the pipelined datapath.(cp2)
- (3) L1 Instruction cache , L1 data cache and arbiter. (cp2)
- (4) Tournament branch predictor. (cp3)
- (5) Pipelined 4-ways cache. (cp3)
- (6) L2 cache. (cp3)
- (7) Victim cache. (cp3)



Tournament branch predictor

Overview:

- Uses two branch history tables
 - Local branch history table (BTB with 6 bits from PC)
 - Global branch history table (BTB with plus a bit sequence which tracks the outcome of the last 4 branches)
- Decides outcome with 2-bit selector which switches between the local and global predictor
 - Similar to the 2-bit branch predictor described in lecture



Challenges:

- Checking for flushing condition
 - Both the prediction and the target address must be correct
- Finding the ideal number of bits to store in the branch prediction table

Tournament branch predictor

Performance Table

	Local: 5 PC bits Global: 4 PC bits, 3 past branch bits	Local: 5 PC bits Global: 4 PC bits, 4 past branch bits	Local: 7 PC bits Global: 5 PC bits, 4 past branch bits
comp1	8860/12528 (70%)	9005/12586 (72%)	10028/12611(80%)
comp2	19802/35414 (56%)	20472/35447 (58%)	25496/35472 (72%)
comp3	3076/4787 (64%)	3172/4754 (67%)	3818/4787 (80%)

Conclusion/Comment: Increasing the number of past prediction bits can help performance, but increasing the number of PC bits is much more effective

Pipelined Cache

Overview:

Using Bram module to achieve next cycle hit, bram enable larger L1 cache with little cost in power and frequencies. Bram module doesn't use FPGA register. When the first tag_out match with mem_address, it is already loading the next tag. So we can do next cycle hit without needing to stall the pipelined CPU.

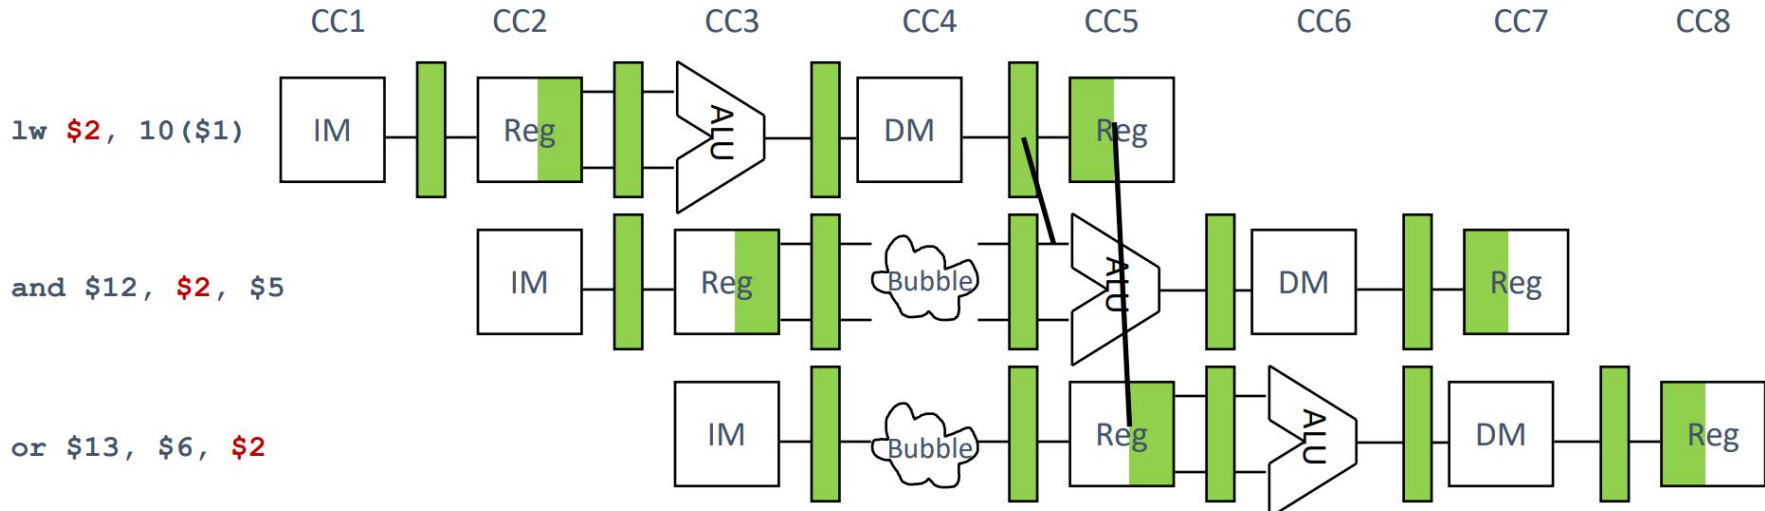
Revising the stage design for I1 icache and dcache.

Changing forwarding path since now we get the data one cycle late or we need to fetch data one cycle early.

Challenges:

- Bram tutorial was given to be very limited. There were a couple of settings that we needed to find out ourselves. The change in forwarding design is needed in both lcache and dcache, namely need to change where instruction data come in and where dcache data come out.
- Implementing bubble states to prevent a data hazard between a load function and a consecutive dependency
- Writing back to memory on dirty evictions

Bubble State Example



Pipelined Cache Performance Table

	DATA CACHE		INSTRUCTION CACHE
	32 Sets, 2 Ways	32 Sets, 4 Ways	32 Sets, 2 Ways
comp1	Misses: 7/3202 (0.21%) Write Backs: 0 Bubbles: 0	Misses: 7/3202 (0.21%) Write Backs: 0 Bubbles: 0	Misses: 25/49431 (.051%)
comp2	Misses: 24/4581 (0.52%) Write Backs: 0 Bubbles: 155	Misses: 24/4581 (0.52%) Write Backs: 0 Bubbles: 155	Misses: 41/108598 (0.038%)
comp3	Misses: 284/11494 (2.5%) Write Backs: 1 Bubbles: 4345	Misses: 283/11374 (2.4%) Write Backs: 1 Bubbles: 4345	Misses: 30/63515 (0.047%)

Conclusion/Comment: For the competition programs, the 4 way D-cache did not significantly improve performance!

L2 cache

Overview:

- We could obtain a L2 cache by removing the bus adaptor of a L1 cache
- The L2 sits in between the arbiter and the physical memory

Challenges:

- Being careful when connecting the ports
- Transferring data between L1 and L2
- Choosing a good size for the L2

L2 Cache Performance Table

	DATA CACHE	
	L2	L1
comp1	Misses: 20/20 Write Backs: 0	Misses: 1/209 (0.21%) Write Backs: 1 Bubbles: 0
comp2	Misses: 20/20 Write Backs: 0	Misses: 24/4581 (0.52%) Write Backs: 0 Bubbles: 155
comp3	Misses: 304/304 Write Backs: 276	Misses: 283/11374 (2.4%) Write Backs: 1 Bubbles: 4345

Conclusion/Comment:

Victim Cache

Overview:

- Create several arrays with length 16 to store valid bit, dirty bit, address(as key), and data.
- Implement lru queue to keep track of lru data.
- Only comes to victim cache on upper cache misses.
- Evicted data whether clean or dirty will be put into victim cache. Victim cache only ever write back when its own evicted data are deemed dirty.
- lcache doesn't do any write back so we only consider dcache/l2cache case.

Challenge: We tried to create one separately so it was first tested on given cache. Given cache was one cycle hit and it ends up creating a lot more trouble than we anticipated, it was extremely difficult to debug using modelsim because every error was created way earlier in the process by wrongly write/writeback data.

Duplicate entry in the victim cache has to be rearrange properly or write back

Victim Cache

Performance on given dcache

```
global_branch_predictor
# Reset Memory
# L1 read-hit in victim:
592, L1 read-miss in victim:
      280, L1 write-hit in victim
:      419, L1 write-miss in
victim:      122
# ** Note: $finish      : /home/s
159/ece411/SnapDragon-411/hvl/s
ource_tb.sv(58)
#      Time: 3713155 ns Iteratio
```

Total misses:1413

Total found in
victim:1011

Total miss in victim:402

Overall Performance

	Comp1 Time (ns)	Comp1 Power (mW)	Comp1 Score
Baseline	720,755.00	448.44	1.68E-10
Snapdragon	332137.9975	607.89	2.22731011E-11
	Comp2 Time (ns)	Comp2 Power (mW)	Comp2 Score
Baseline	4,521,285.00	430.48	3.98E-08
Snapdragon	753166.1535	978.48	4.18046256E-10
	Comp3 Time (ns)	Comp3 Power (mW)	Comp3 Score
Baseline	3,639,265.00	425.27	2.05E-08
Snapdragon	421858.0535	855.98	6.42632449E-11

Geometric Mean $= (\text{Comp1Score} * \text{Comp2Score} * \text{Comp3Score})^{(1/3)} = 8.42666817 \times 10^{-11}$

Q&A

Thanks!

